

What changes from Red Shell

Most of the Red Shell foundation carries over. These four rules moved, and breaking any one of them now fails a task outright.

1 One prompt, one answer

The task type is always **single turn**: Model A gets the initial prompt and no follow ups. A Model A conversation with more than one user turn **fails the task**.

- **No milestones, no deferred files, no revision turns.** Every requirement goes into the one prompt.
- Leg B still runs the same prompt word for word, steered with **intent level hints**; a hint that supplies a value the model then uses fails the task.

Guidelines · 1.1.2 Rest of the Parameters · QC spec · Task Type - Single-Turn Structure

2 Model A has to fail at least 30%

Genuine failures worth **at least 30% of the final rubric score**, with real impact on the deliverable, never forced or artificially scored. **50% or more is preferred**.

Model A scores 70% or less

Its failures clear the 30% floor

Above 70%

Fails as a Trivial Task

Guidelines · Hard Client Requirements for Every Task · QC spec · Rubric Criteria - All Criteria Scoring

3 Two automatic fails in the rubric

Completion · 80% or more

What was delivered: the artifact, the state change, the final message

Process ·

20% max

How it got there

- ✗ **Process over the cap.** Count the objective block and round down: 12 criteria allow 2. Zero is preferred.
- ✗ **Any existence check.** A criterion that only verifies a file, section, column or record is there.

Guidelines · 5.1 Rubric Fundamentals · 5.6 How the Client Defines Rubric Issue Severity

4 Every assigned parameter is binding

All **fourteen** parameters, from the scenario and the assigned tools to the Execution target, are followed exactly. **Omitting, substituting or drifting from one makes the task invalid and subject to automatic rejection.**

- Claiming the universe cannot support the scenario when it can **fails the task automatically**: explore before you deviate.
- **11 use cases and 68 subcategories** replace the old 6 and 16.

Guidelines · Hard Client Requirements for Every Task · 1.1 Task Parameters & Execution Rules

ALSO NEW At least 3 multimodal inputs, most images needing real visual interpretation · **Every literal** in a criterion copied from its source, subjective block included

Common Errors

The mistakes that keep costing people their tasks on the single turn project, and what to do instead.

Every error in this deck cites its source. Each one comes with a direct reference to the exact section of [the Green Shell Guidelines](#) or [the QC spec](#) that sets the correct standard, so you can check it yourself. **Each example number is a link** to the real Green Shell task where it happened.

Every example in this deck comes from a **real task submitted to Green Shell**. Many of these errors break the most basic standards of the project, and correcting them usually takes very little extra effort.

These are not minor issues. Repeating these mistakes can directly affect your access to tasks and your ability to remain on the project.

You can go through this course superficially and risk making the same mistakes again, or use it to understand exactly what is going wrong and what needs to change. Going forward, these standards are expected to be applied consistently. **If the same fundamental issues appear again after this guidance, continued access to tasks and to the project should not be expected.**

HOW TO USE THIS DECK

Fix order is build order

The seven steps of a Green Shell task, and the sections whose mistakes are made at each. The earlier the step, the more a fix costs.

STEP	WHAT YOU ARE DOING	SECTIONS WHOSE MISTAKES ARE MADE HERE
1-2	Plan and set up · universe, inputs and the prompt BEFORE YOU RUN Guidelines · 1 Planning the Agent Task · 2 Setting Up the Run	1 The Foundations Every Task Must Get Right · the parameters and the inputs 2 A Rubric That Grades the Outcome · every ask the prompt makes
3-4	Run Model A · one turn, no follow ups LEG A · MODEL A Guidelines · 3 Task Runner - Leg A	5 Closing the Task · the trajectory and the artifacts, downloaded
5	Write the objective rubric, rate Model A LEG A · MODEL A Guidelines · 5 Objective Rubrics	2 A Rubric That Grades the Outcome 3 Each Check Once, at the Weight It Earns
6	Reach the golden · hint Model B until it passes LEG B · MODEL B Guidelines · 6 Golden Solution - Leg B	5 Closing the Task · every hint, and the golden itself
7	Write the subjective block, rate Model A and the golden LEG B · MODEL B Guidelines · 7 Subjective Rubric Block	4 The Subjective Block
	Before you submit	5 Closing the Task · the preferred run

The Foundations Every Task Must Get Right

Two decisions made while the task is still on paper. Fixing either after Model A has run means reworking the task and running Leg A again, because the prompt, the inputs and every criterion inherit them.

Every assigned parameter, built in

The fourteen parameters are binding client requirements. Build the scenario on the assigned tools, put the deliverable where the Execution target says, and keep the scenario's user, goal and deliverable.

- Assigned tools** · name the step each one is necessary for; the correct final state depends on them
- Execution target** · stated in the prompt and graded, such as drafts left unsent in the mail client
- Scenario** · fill in only what it leaves open, and deviate only for a real universe limit or the failure threshold
- Universe** · explore before you drift: misadherence from thin exploration fails automatically

Why Omitting, substituting or drifting from any assigned parameter makes the task invalid and subject to automatic rejection.

SOURCE
Guidelines · Hard Client Requirements for Every Task
Guidelines · 1.1.2 Rest of the Parameters

SEE EXAMPLES [1](#) AND [2](#)

EXPLICIT CLIENT REQUIREMENT

Images worth looking at, not only worth reading

Most images have to need genuine visual interpretation: photographs, diagrams, charts, maps, real world objects. An image the model only transcribes is a document in picture form.

- At least 3 multimodal inputs** · in the Model A conversation, a floor and not a target
- Every assigned modality** · used, each input carrying part of the answer
- Ask of every image** · what does the model have to do with it? If the answer is read the words, it is a document

Why The client expects real image diversity: most of the images used as multimodal context must not be primarily text based.

SOURCE
Guidelines · Hard Client Requirements for Every Task · Multimodal
Guidelines · 1.2.2 Select the Multimodal Inputs

SEE EXAMPLE [3](#)

NEW IN GREEN SHELL

A Rubric That Grades the Outcome

The client grades completion: did the agent produce the artifact, follow the instructions and deliver the outcome the user asked for? Three standards keep a rubric on what was delivered.

Grade it where it lands

A reasoning decision is graded in the value, classification or decision written into the deliverable, not as a step in the trajectory. Keep a process criterion only where no deliverable can show the failure.

Completion, 80% or more · Task Completion and Instruction
Following on the artifact, the state change or the message

Process, 20% at most · Tool Use, Agent Behavior, any Trajectory target, narration

Count the objective block · and round down: 12 criteria allow 2, 28 allow 5

Why Above 20% the task fails automatically; an avoidable process criterion inside the cap still costs points.

SOURCE
Guidelines · 5.1 Rubric Fundamentals
QC spec · Rubric Criteria - 80/20 Outcome Split (Process Over Cap)

SEE EXAMPLES [4](#) AND [5](#)

Never grade existence

A criterion that passes as soon as a file, section, column or record is there grades nothing. Existence is already implied by every criterion on the content.

The test · could it pass on a file full of wrong content?

Disguised forms · a section title, a list of columns, a chart produced, an email sent, with no content graded

A filename alone · the same check in another form, in either block

Why A single existence check fails the task automatically, whatever the rest of the block looks like.

SOURCE
Guidelines · 5.1 Rubric Fundamentals
QC spec · Rubric Criteria - Existence Check

SEE EXAMPLE [6](#)

Every ask graded

Walk the prompt once and tick off every ask: anything you can tick needs a criterion. Where the deliverable ends up is an ask too, whenever the prompt says it.

Number the asks · and write a criterion number next to each one

The other direction · every criterion traces back to a line of the prompt or the inputs

Why A missing criterion for a critical requirement is Major: the grade comes out wrong, and Major issues above 10% fail the task.

SOURCE
Guidelines · 5.7 Final Rubric Checklist
Guidelines · 5.6 How the Client Defines Rubric Issue Severity

SEE EXAMPLE [7](#)

Each Check Once, at the Weight It Earns

A rubric can grade the right things and still score them wrong: counted twice, weighted on the restatement, or tagged as the wrong kind of check. Each survives a quick read; none survives a count.

One outcome, one criterion

Each criterion captures one grading signal and is decided by its own evidence. A criterion restating what others grade, often a long Trajectory line that bundles the answer key, counts one slip twice.

- Parts that could pass separately** · split them
- Evidence another criterion decides** · delete the restatement

Why Overlapping or Redundant is a Moderate issue, and two to four issues of any severity can reject a task.

SOURCE
Guidelines · 5.1 Rubric Fundamentals
Guidelines · 5.6 How the Client Defines Rubric Issue Severity

SEE EXAMPLE 8 ↗

Weight is difficulty, never importance

Count the axes a criterion exercises: source coordination, reasoning depth, modality, discovery effort. Three or more is +5, exactly two is +3, one or none is +1.

- +5** · evidence integrated across modalities, sources or reasoning steps
- +1** · a direct copy or lookup from one source
- Never on a restatement** · the weight sits on the reasoning act

Why A finding that took a search and an inference priced at +1, with its restatement at +5, lets the route decide the score.

SOURCE
Guidelines · 5.2 Rubric Weights

SEE EXAMPLE 9 ↗

The category says what kind of check it is

Target is mechanical: where the grader looks. Category is the kind of check, and Task Completion is the default: anything the model can get wrong.

- Task Completion** · the thing the user came for; a model can be wrong at it
- Instruction Following** · a stipulation; a model can only disobey it
- Agent Behavior** · always on Trajectory, counted in the 20%

Why Category and target decide how the rubric data is structured and read downstream by the client.

SOURCE
Guidelines · 5.4 Category & Evaluation Target

SEE EXAMPLE 10 ↗

The Subjective Block

The render side nobody asked for, judged as the user would: ten or more criteria, each one element and one visible property, all of them Task Completion.

Ten or more criteria, from a render rich enough

Compare the Model A and Model B artifacts side by side as the user, and write a criterion for each visible property that improves or worsens the experience. Expect ten or more; never pad to reach them.

Cover the render · layout, hierarchy, spacing, readability, structure, consistency, register

One element, one visible property · never a vague judgment such as looks professional

Weight is impact on the user · not how hard the property was to achieve

Why Subjective rubrics play a core role in the data the client expects, and an ungraded part of the render scores both models the same.

SOURCE

Guidelines · 7.1 Writing Subjective Rubrics

Guidelines · 7.2 Subjective Rubric Weights

SEE EXAMPLES [11](#) AND [12](#)

Task Completion, on the render

Task Completion is the only valid category in the subjective block, on the Final Answer Artifact almost always. Anything the prompt asked for is objective only.

Instruction Following · excluded by design

Agent Behavior, Tool Use, Safety · judged from the trajectory, not the render

Never a value, never existence · two value checks fail the task, and so does a single existence check

Why A requirement the prompt states belongs in the objective block, and a criterion with the wrong category reaches the client mislabelled.

SOURCE

Guidelines · 7.4 Categories and Evaluation Targets for Subjective Rubrics

QC spec · Rubric Criteria - Subjective Block Scope

SEE EXAMPLE [13](#)

Closing the Task

Leg B runs the same prompt word for word and steers Model B to the ideal response. QC reads these straight from the runs, so they are quick to check and costly to miss.

A hint points at the area, never at the answer

Say what is still missing, from the user's point of view. Never what to write: no value, position or conclusion the model should find in the inputs or the universe.

LEAK

"Add the corrected price to the audit."

HINT

"I still cannot see anything from the studio walkthrough reflected in the audit."

- ✖

Fails the task · a hint supplies a value the model uses without going back to the source
- !

Costs points · over steering, such as naming the exact file to open

Why The golden is the reference the subjective block and the pairwise comparison are judged against.

SOURCE
QC spec · Golden Solution - Hint Leak (Leg B)
Guidelines · 6.2 Steering the Run & Hinting

The golden ships finished, grounded and complete

The same deliverable in its best form: every file the prompt names, spelled verbatim, every factual claim grounded in the inputs, artifacts only.

- ✖

Fails the task · a file the prompt names is missing
- !

Costs points · an inaccuracy, an ungrounded claim, narration of how it was made
- ✓

It passes every objective criterion · anything it fails is a broken criterion, not a broken golden

Why A weak or defective golden distorts the subjective evaluation that is scored against it.

SOURCE
QC spec · Golden Solution - Artifact Completeness
Guidelines · 5.7 Final Rubric Checklist

Close the runs out

Download the trajectory and the final artifacts as soon as a run completes, and rate Model A from the trajectory, the artifacts and the GTFA together.

- ✖

Fails the task · a trajectory is missing
- !

Costs points · no preferred run marked when more than one exists
- ✓

Every failure justified · why the rubric is correct, where Model A failed, why it is necessary

Why Without a preferred run there is no reliable way to tell which run is the right one.

SOURCE
Guidelines · 3. Task Runner - Leg A
Guidelines · 5.5 Rating Model A and Justifying the Failures
QC spec · Trajectory - Completeness · Golden/Preferred Run Selection

The checks, section by section

If you read one slide twice, read this one. One check per standard in this deck, in the order you build the task.

EVERY TASK

- 1 One prompt to Model A, no follow ups
- 2 Model A genuinely fails at least 30% of the rubric
- 3 Every literal in a criterion copied from its source

3. EACH CHECK ONCE, AT THE WEIGHT IT EARNS

- 1 No two criteria are decided by the same evidence
- 2 Each weight matches the axes its criterion exercises, never on a restatement
- 3 Task Completion for anything a model can get wrong; Agent Behavior on Trajectory

1. THE FOUNDATIONS EVERY TASK MUST GET RIGHT

- 1 Each assigned tool is necessary to a step; the Execution target is in the prompt and graded
- 2 The scenario's user, goal and deliverable are intact
- 3 At least 3 multimodal inputs, most images needing to be looked at, not read

4. THE SUBJECTIVE BLOCK

- 1 Ten or more subjective criteria, each one element and one visible property
- 2 Every subjective criterion is Task Completion, graded on the render

2. A RUBRIC THAT GRADES THE OUTCOME


- 1 Process criteria at most 20% of the objective block, each one impossible to grade on the deliverable
- 2 No criterion could pass on a file full of wrong content
- 3 Every ask in the prompt has a criterion, and every criterion traces to the prompt or the inputs

5. CLOSING THE TASK

- 1 Every Leg B hint stays at intent level
- 2 The golden ships every named file, grounded, and passes every objective criterion
- 3 Trajectories and artifacts downloaded; the preferred run marked

What happens from here is simple

You have seen the mistakes that keep costing tasks and project access. Nobody is expected to be perfect, but once an issue has been identified and explained, it is expected to be corrected, not repeated.



These are not hidden expectations anymore: the standards, the real examples and how to meet them are all in this course.

The goal is not to catch you making mistakes. It is to make sure the same preventable mistakes stop happening.

Thank you for your engagement and cooperation.